

# Text Mining

## Introduction to Text Mining

Academic Year 2024-2025

Bruno Jardim    Rita Oliveira

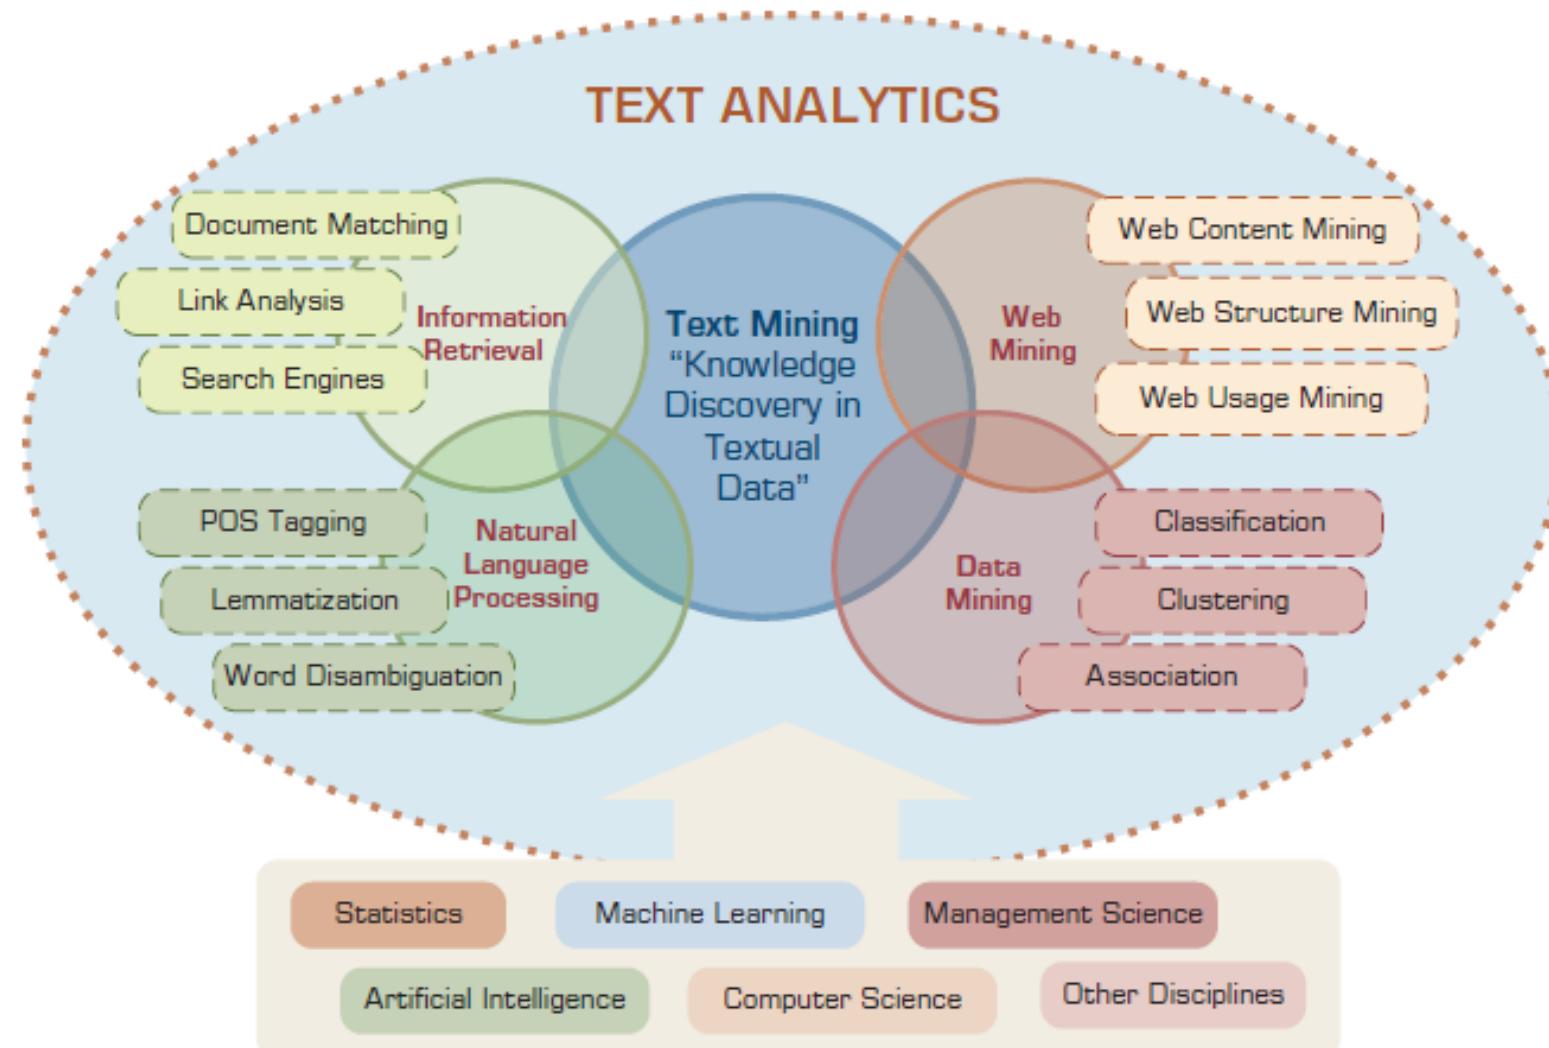
[bjardim@novaims.unl.pt](mailto:bjardim@novaims.unl.pt)    [roliveira@novaims.unl.pt](mailto:roliveira@novaims.unl.pt)

# Lecture Plan

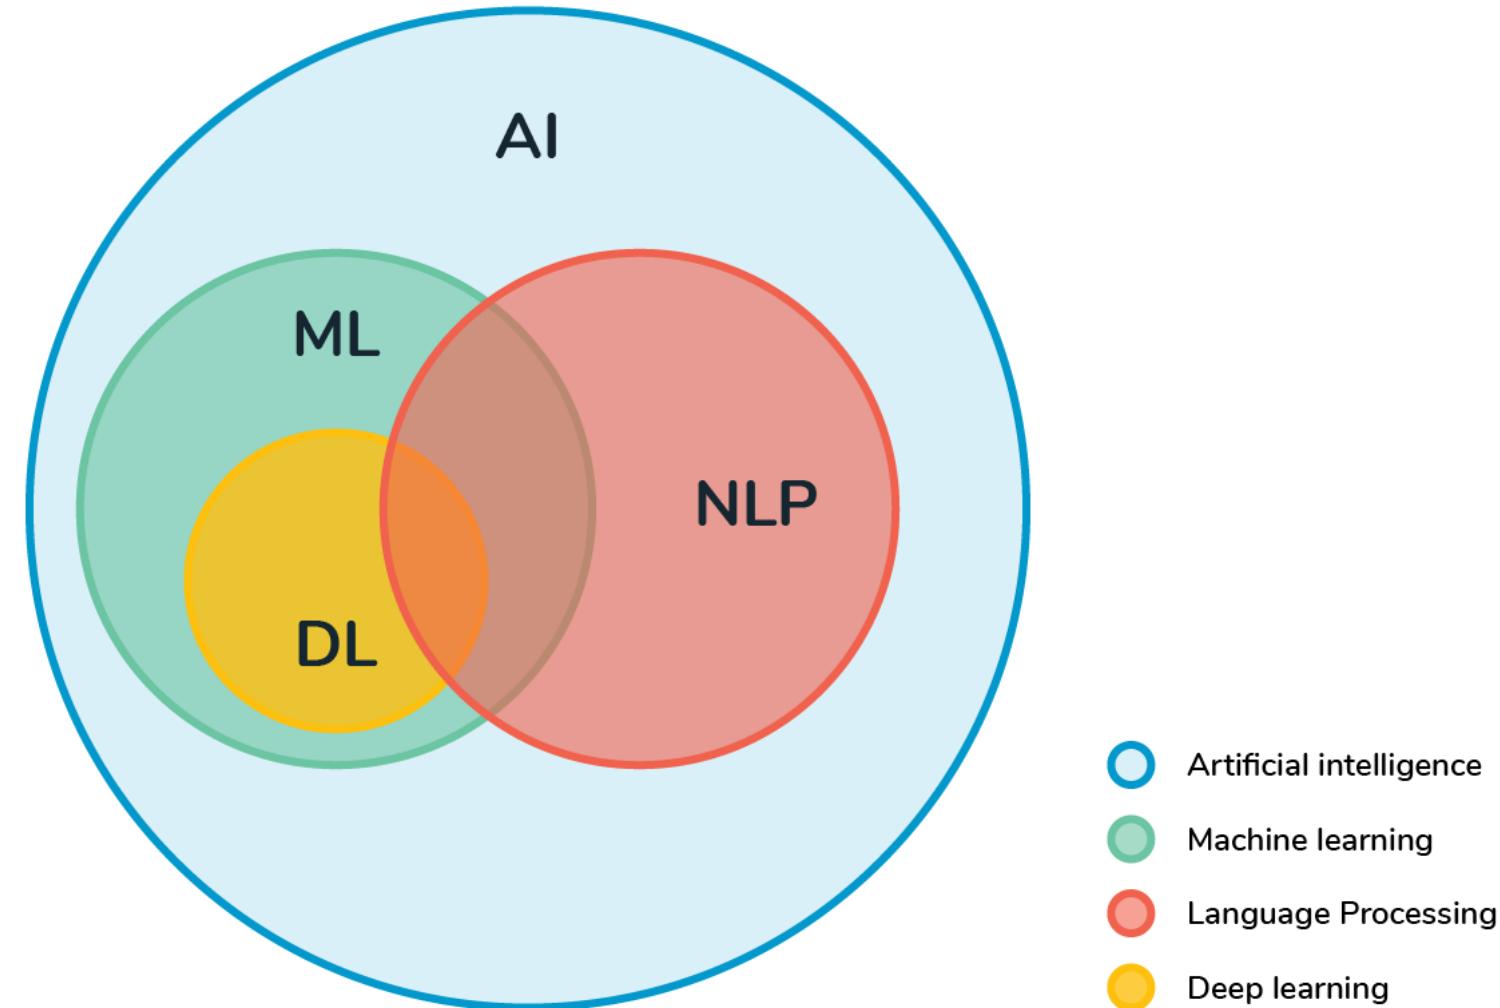
1. Introduction to Text Mining
2. Natural Language Processing (NLP) Pipeline
3. Text Preprocessing

# 1. Introduction to Text Mining

# Introduction to Text Mining



# Introduction to Text Mining

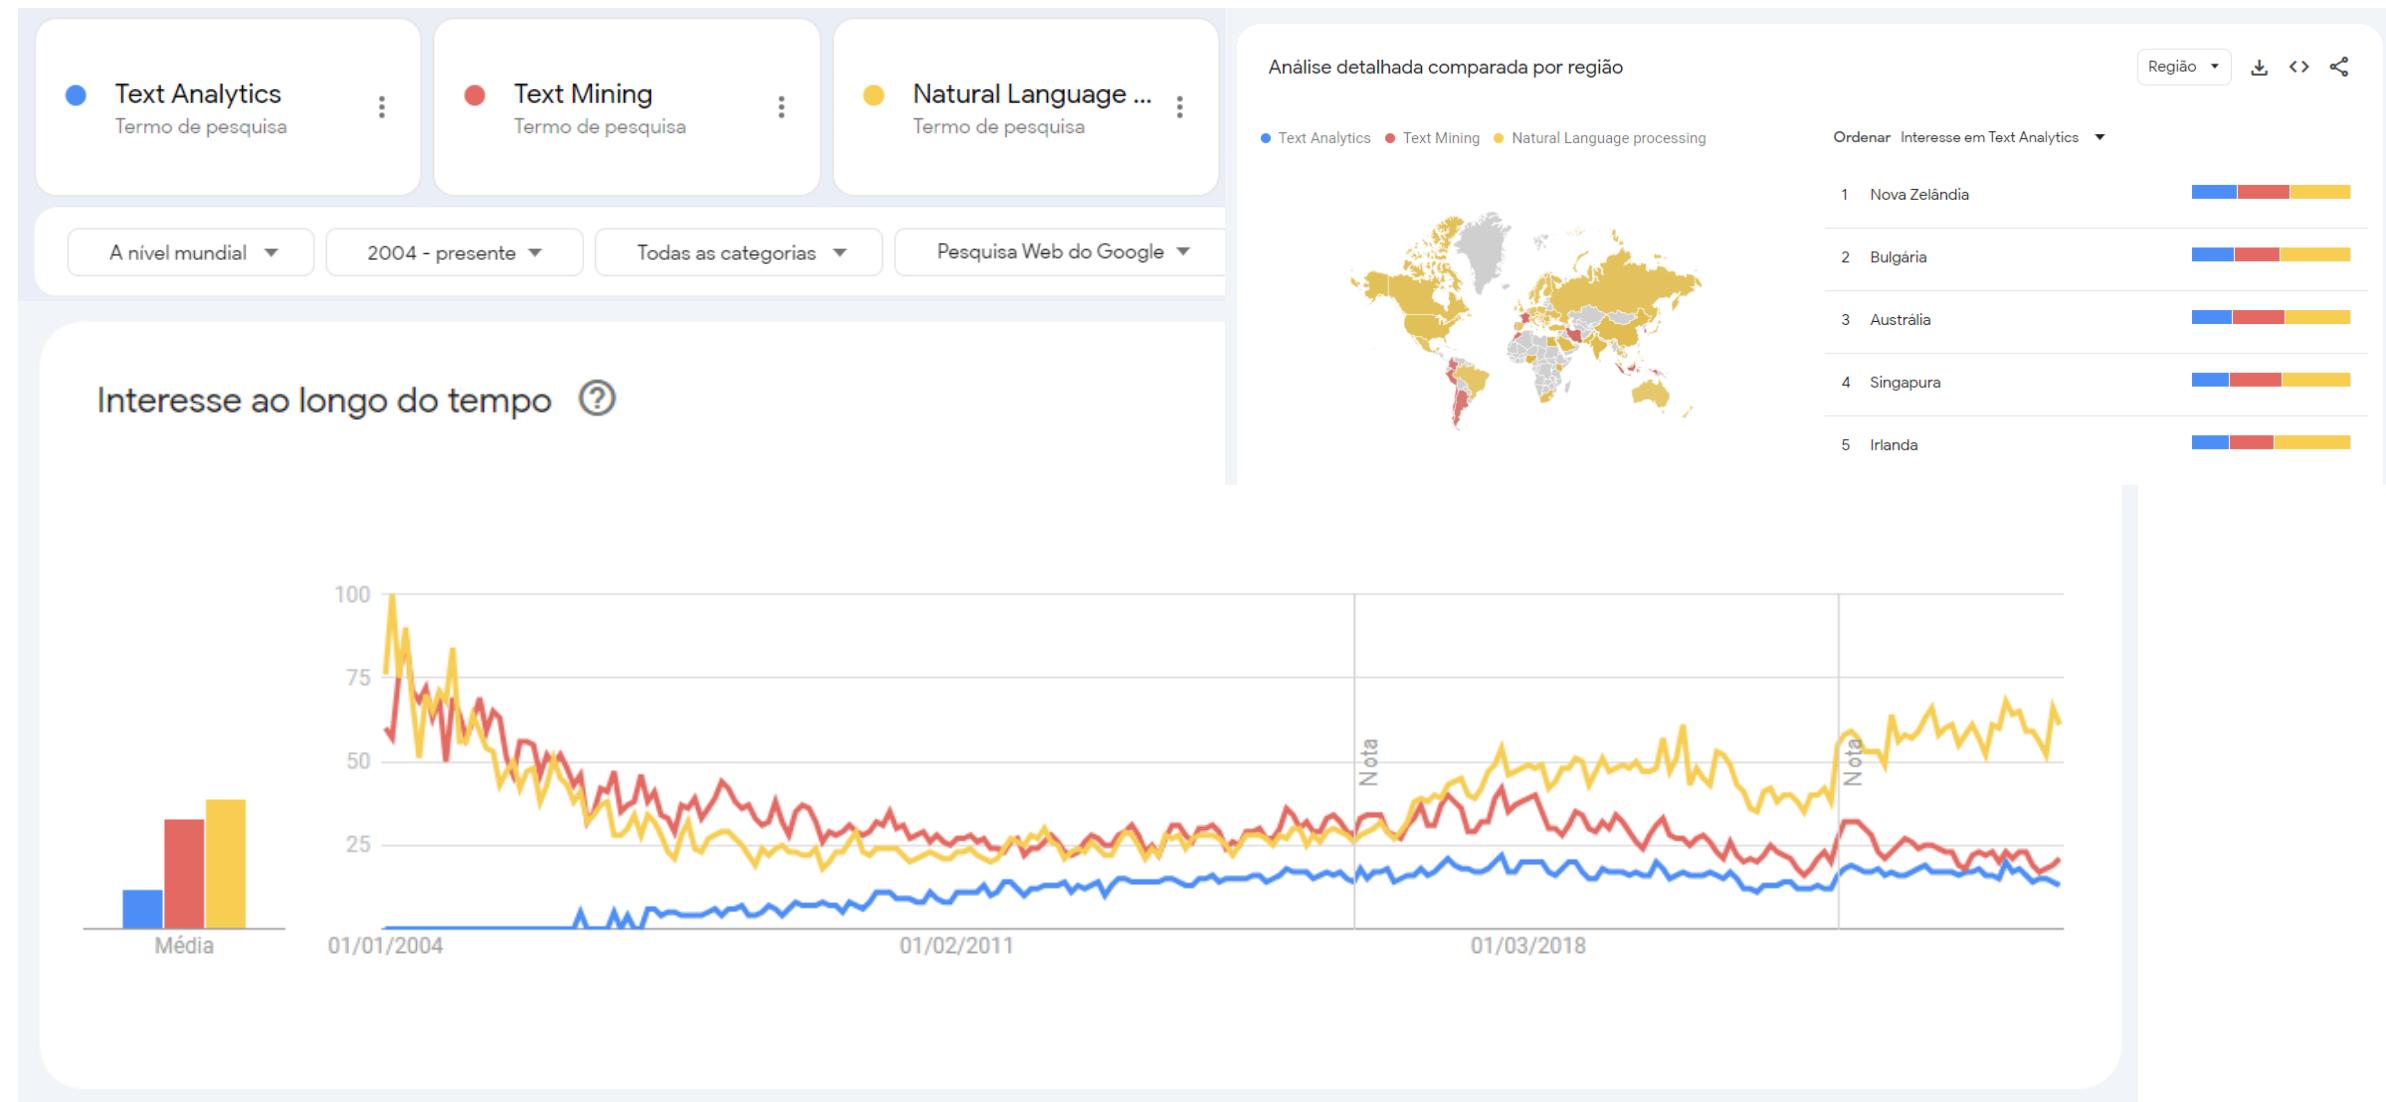


# Introduction to Text Mining

**Text mining** (also referred to as ***text analytics***) is an **artificial intelligence (AI) technology that uses natural language processing (NLP)** to transform the free (unstructured) text in documents and databases into normalized, structured data suitable for analysis or to drive machine learning (ML) algorithms.

**Natural language processing (NLP)** is an area of computer science and artificial intelligence that is concerned with the interaction between computers and humans in natural language. The **ultimate goal of NLP is to enable computers to understand language as well as we do**.

# Introduction to Text Mining



# Applications

1. Machine translation
2. Dialog systems (chatbots)
3. Search Engines
4. Predictive Keyboards and auto correctors

...

**Everything that deals with text!!!**



# Applications

## Natural Language (NL):

1. Grammatical system, with its **own rules**, used by people to communicate
2. Natural **evolution** due to people communication
  - ✓ New words everyday: “bue”, “fixe”, “ups”, “lol”

## Examples:

1. English
2. Portuguese

...

Python??

# Challenges of NL - Variability

## Definition:

Different **sentences can have the same meaning**; thus, we can say the same thing in different ways. Two sentences with the same meaning are called **paraphrases**.

## Examples:

“The president greets the press in Lisbon”

VS

“Marcelo speaks to the media in Campolide”

“He has tons of stuff to throw away”

VS

“He needs to get rid of a lot of junk.”

# Challenges of NL - Ambiguity

## Definition:

A **single sentence can have different meanings**

## Dealing with it:

The only way to deal with ambiguity is through context!



# Challenges of NL - Generalization

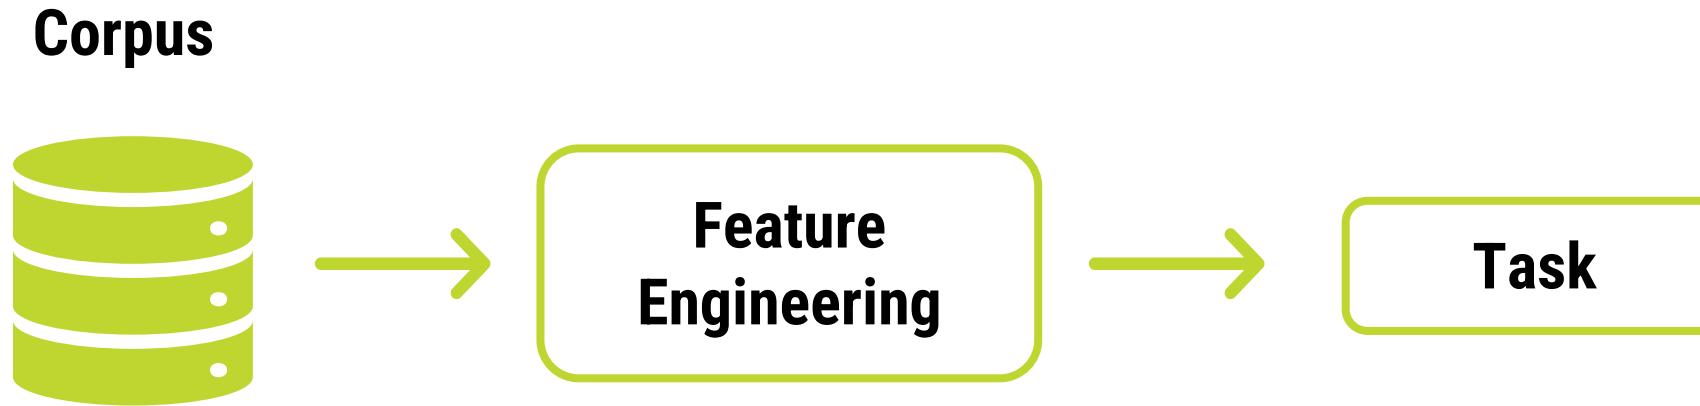
Often an **NLP system is trained with a corpus** in a specific domain, but it is **used in a different domain**.

Thus, the system is confronted with input that he has never seen before, either because those inputs are Out-of-Domain (OOD) or because they contain words Out-of-Vocabulary (OOV).

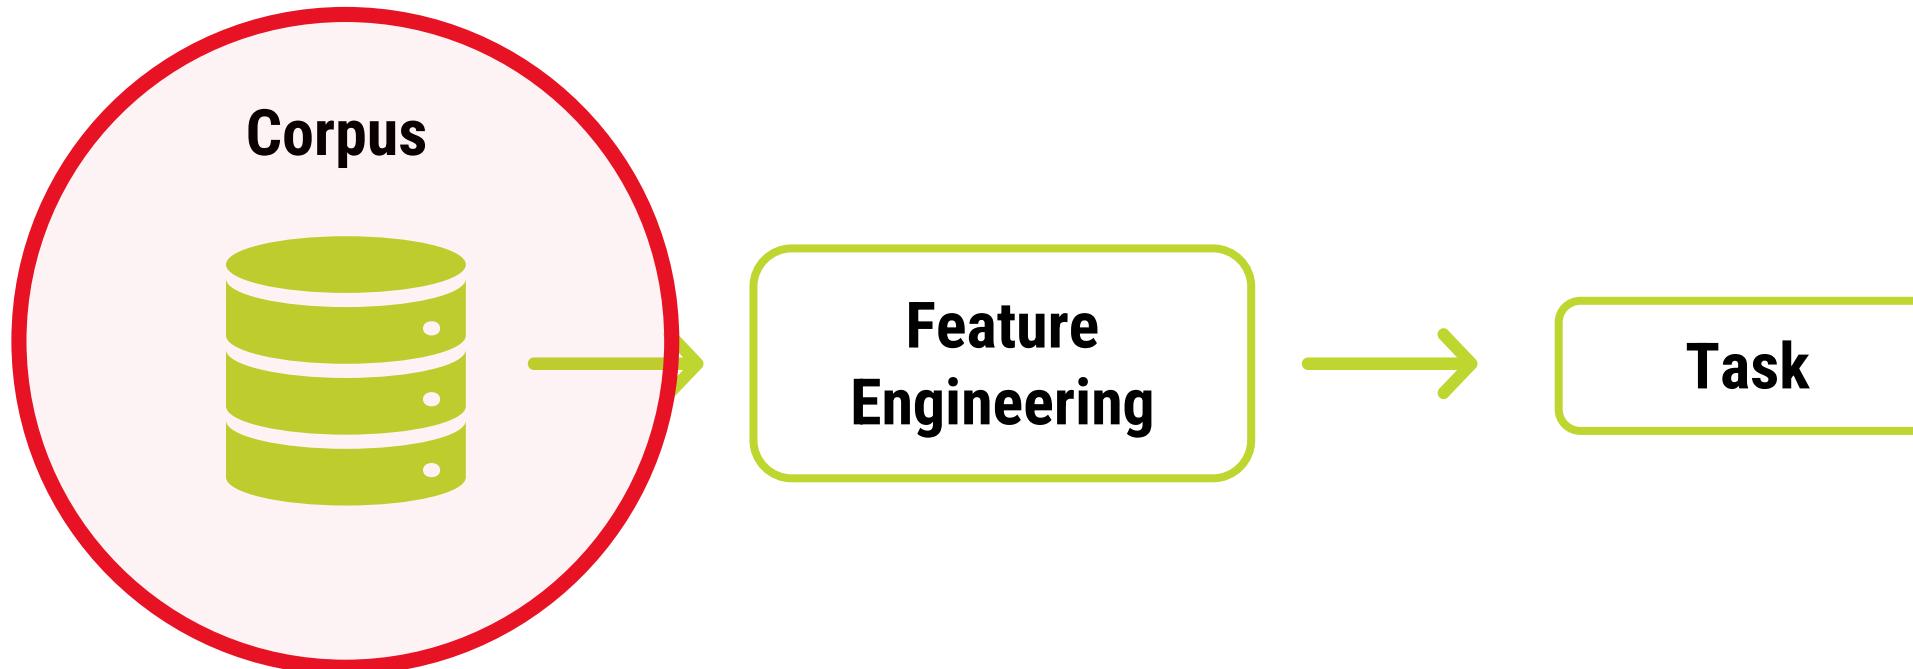


## 2. Natural Language Processing

# NLP Pipeline



# NLP Pipeline



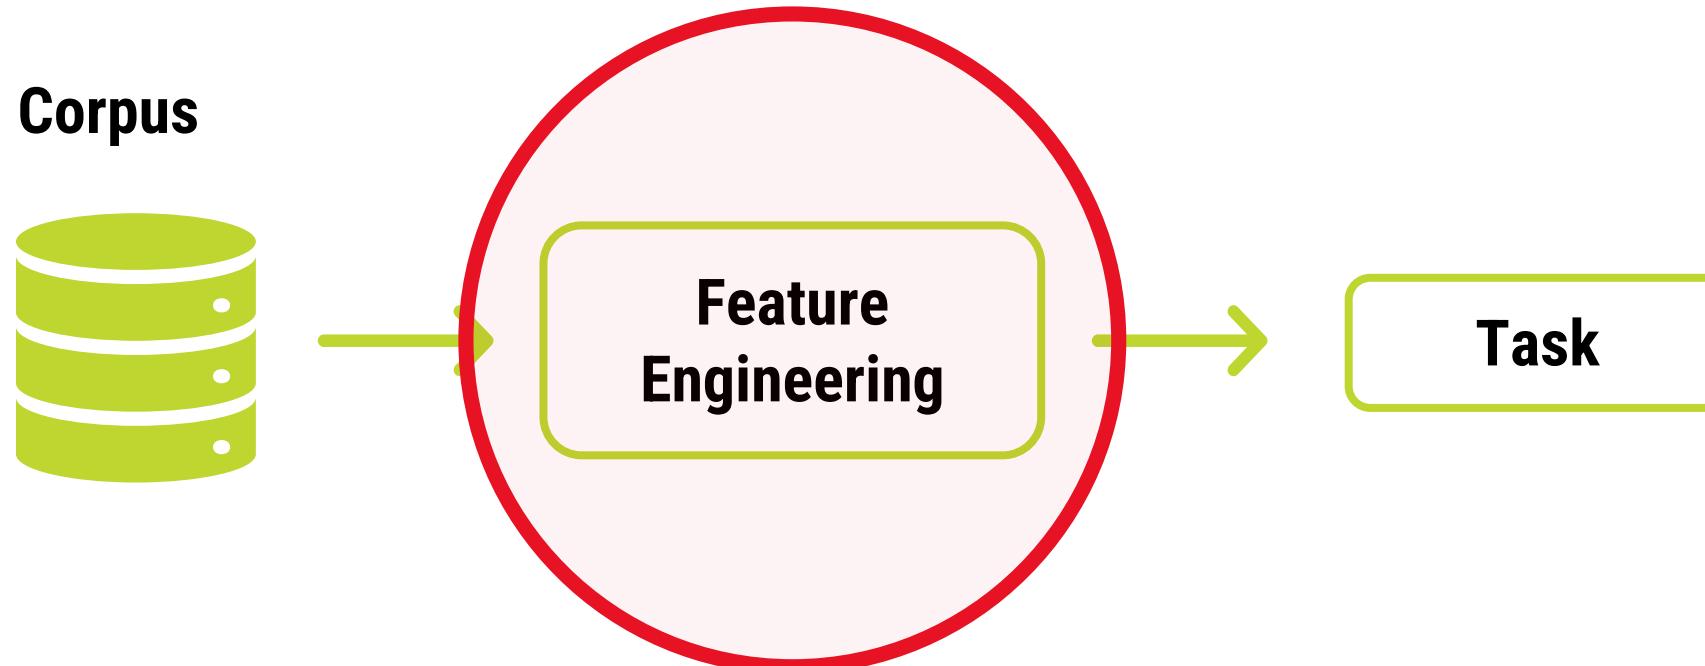
# Corpus

- A **Corpus** is a collection of **text** organized into datasets.
- A corpus can be made up of everything from **news**, **recipes**, **wikipedia pages**, **academic papers**, to **novels**, **movie** and **tv scripts**, and **social media posts** like **tweets**.
- A collection of Corpus is called **Corpora**.

# Corpus

- The first step is to split a corpus into **Train/Validation/Test** or via **K-fold cross validation** (other methods can be considered).
- For **Train/Validation/Test**
  - 80%/10%/10% spill for small corpus (<10k samples )
  - More train percentage for bigger corpus.
  - Keep the original always!
  - Your split must be reproducible

# Feature Engineering



# Feature Engineering

**How can we represent text?**



# Feature Engineering

## Bag-of-Words

- Each word is a feature.
- Our **feature space is defined by our vocabulary**
- Documents/pieces of text will be represented as **sparse vectors**.

the dog is on the table

0 0 1 1 0 1 1 1  
are cat dog is now on table the

# Feature Engineering

## Bag-of-Words Example

Text 1 – “I love Paris”

Text 2 – “ I live in France”

|        | France | I | in | live | love | Paris |
|--------|--------|---|----|------|------|-------|
| Text 1 | 0      | 1 | 0  | 0    | 1    | 1     |
| Text 2 | 1      | 1 | 1  | 1    | 0    | 0     |

# Feature Engineering

## Bag-of-Words Example

Text 1 – “I love Paris”

Text 2 – “I live in France”

# Feature Engineering

## Bag-of-Words Example

Text 1 – “I love Paris”

Text 2 – I live in France”

|        | France | I | in | live | love | Paris |
|--------|--------|---|----|------|------|-------|
| Text 1 |        |   |    |      |      |       |
| Text 2 |        |   |    |      |      |       |

# Feature Engineering

## Bag-of-Words Example

Text 1 – “I love Paris”

Text 2 – “I live in **France**”

|        | France | I | in | live | love | Paris |
|--------|--------|---|----|------|------|-------|
| Text 1 | 0      |   |    |      |      |       |
| Text 2 | 1      |   |    |      |      |       |

# Feature Engineering

## Bag-of-Words Example

Text 1 – “I love Paris”

Text 2 – “I live in France”

|        | France | I | in | live | love | Paris |
|--------|--------|---|----|------|------|-------|
| Text 1 | 0      | 1 |    |      |      |       |
| Text 2 | 1      | 1 |    |      |      |       |

# Feature Engineering

## Bag-of-Words Example

Text 1 – “I love Paris”

Text 2 – “I live **in** France”

|        | France | I | in | live | love | Paris |
|--------|--------|---|----|------|------|-------|
| Text 1 | 0      | 1 | 0  |      |      |       |
| Text 2 | 1      | 1 | 1  |      |      |       |

# Feature Engineering

## Bag-of-Words Example

Text 1 – “I love Paris”

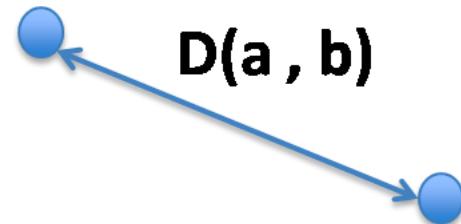
Text 2 – “I live in France”

|        | France | I | in | live | love | Paris |
|--------|--------|---|----|------|------|-------|
| Text 1 | 0      | 1 | 0  | 0    | 1    | 1     |
| Text 2 | 1      | 1 | 1  | 1    | 0    | 0     |

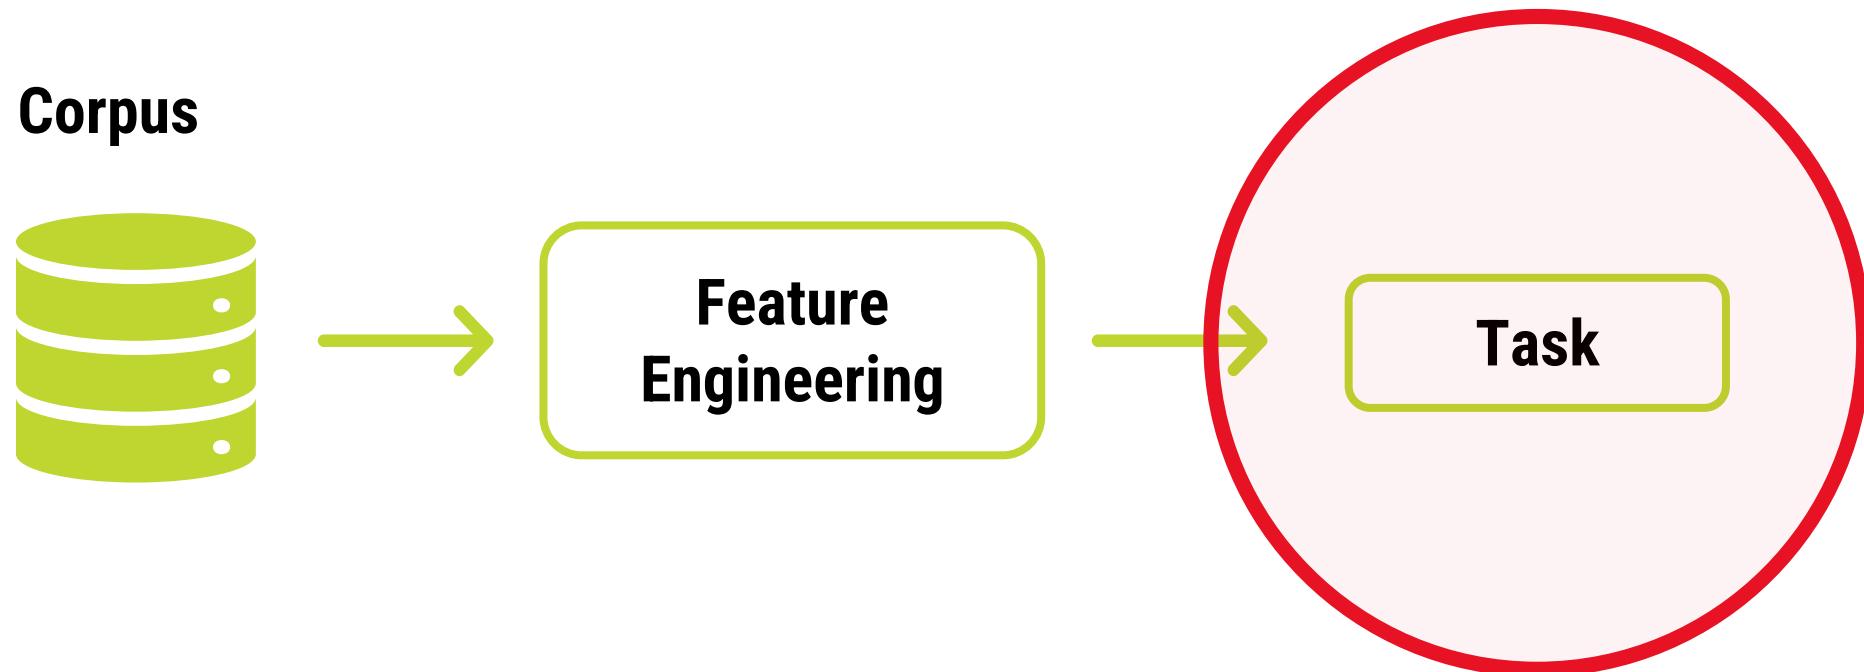
# Feature Engineering

- Given 2 documents we can transform them into sparse vectors and compare them with simple distance metrics.
- If they share a lot of words, they will be close to each other.

$$D(a, b) = \sqrt{\sum_{i=1}^n (b_i - a_i)^2}$$



# NLP Tasks



# NLP Tasks

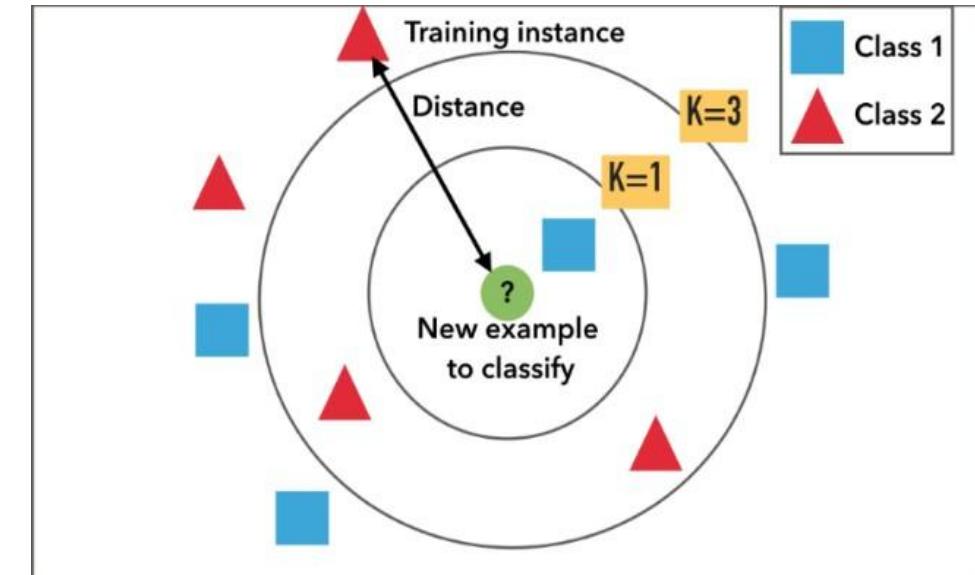
- Keyword Extraction
- Text Similarity
- Document Classification
- Sentiment Analysis
- Topic Modelling
- Information Retrieval
- Q&A
- Text Generation
- (...)

# NLP Tasks

## K-Nearest-Neighbour: (our baseline)

The simplest classification algorithm!

We take our new document  $x$ , and we represent  $x$  in the **same feature space of our training documents**.



Then we **compare  $x$  with all training documents** and we **give  $x$  the label of the closest document**.

# Example of an applied case (cont.)

{you have the original academic article in <https://doi.org/10.1016/j.cstp.2022.07.011>}

## “Predicting illegal parking classes from tickets’ description”:

- Covers 3-year period (2017-2020)
- 89,126 tickets included in analysis
- 7 classes were considered (crosswalk, sidewalk, conditions access, disabled, reserved, others and unknown).

# Example of an applied case (cont.)

## Descriptive features of illegal parking occurrences in Lisbon

| Feature     | Description   |
|-------------|---|
| Datetime    | Date and time of illegal parking occurrence                                     |
| Description | Details and explanation of the illegality as written by the responsible officer |
| Latitude    | Latitude of illegal parking occurrence  |
| Longitude   | Longitude illegal parking occurrence  |
| Address     | Address of illegal parking occurrence   |

# Example of an applied case (cont.)

## Descriptive features of illegal parking occurrences in Lisbon

| Feature     | Description   |
|-------------|---|
| Datetime    | Date and time of illegal parking occurrence                                     |
| Description | Details and explanation of the illegality as written by the responsible officer |
| Latitude    | Latitude of illegal parking occurrence  |
| Longitude   | Longitude illegal parking occurrence  |
| Address     | Address of illegal parking occurrence   |

Unstructured and not systematized!

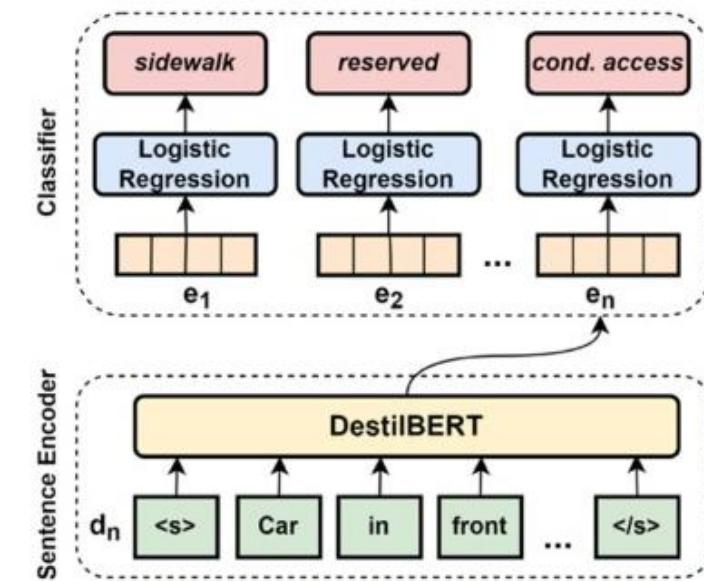
Examples:

- "Parking ticket issued for 'conditioning access' - obstructing driveway entrance."
- "Violation for car parked on sidewalk - hindering pedestrian passage."
- "Ticket given for impeding safe crossing for pedestrians."

# Example of an applied case (cont.)

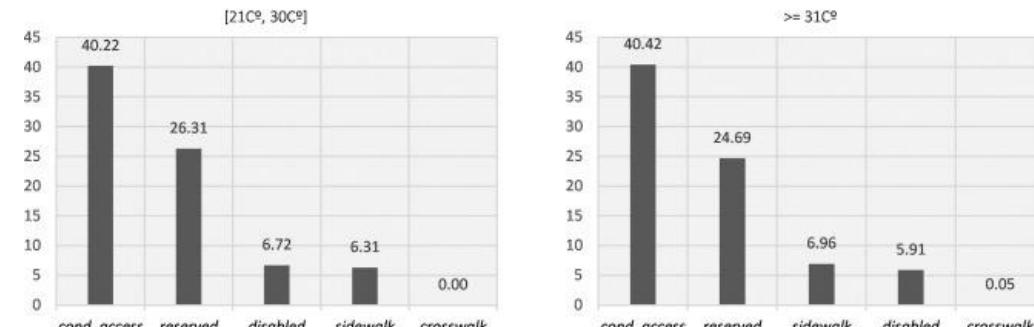
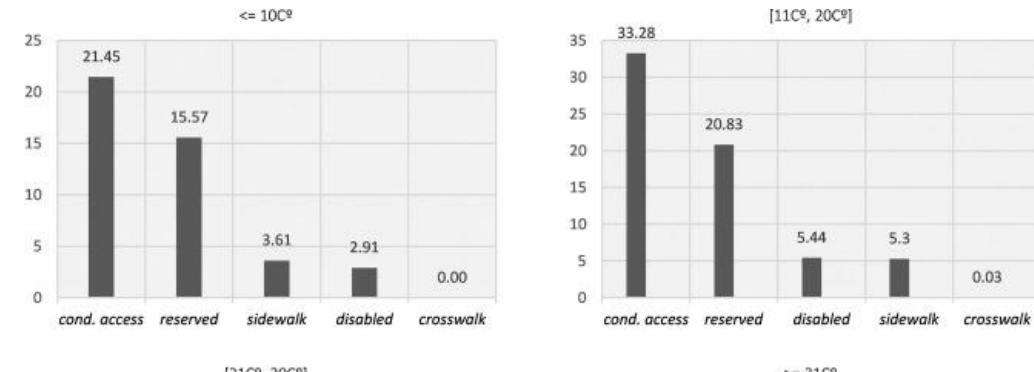
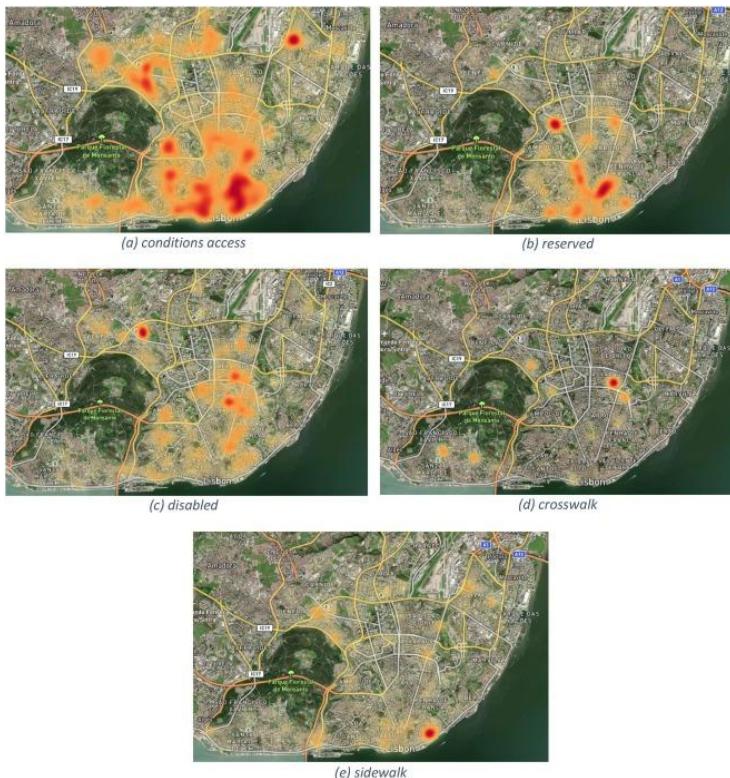
## Text Classification model using NLP state-of-the art models

- NLP was used to **classify the descriptions into pre-defined classes**.
- First, to turn the description text into machine-readable input, it is required to transform it into a vectorial representation (embeddings).
- So, a sentence encoder (DestilBERT - we will learn how to use it later 😊) was used to transform each sentence into a vector.
- A logistic regression receives the input vectors and classifies them accordingly.



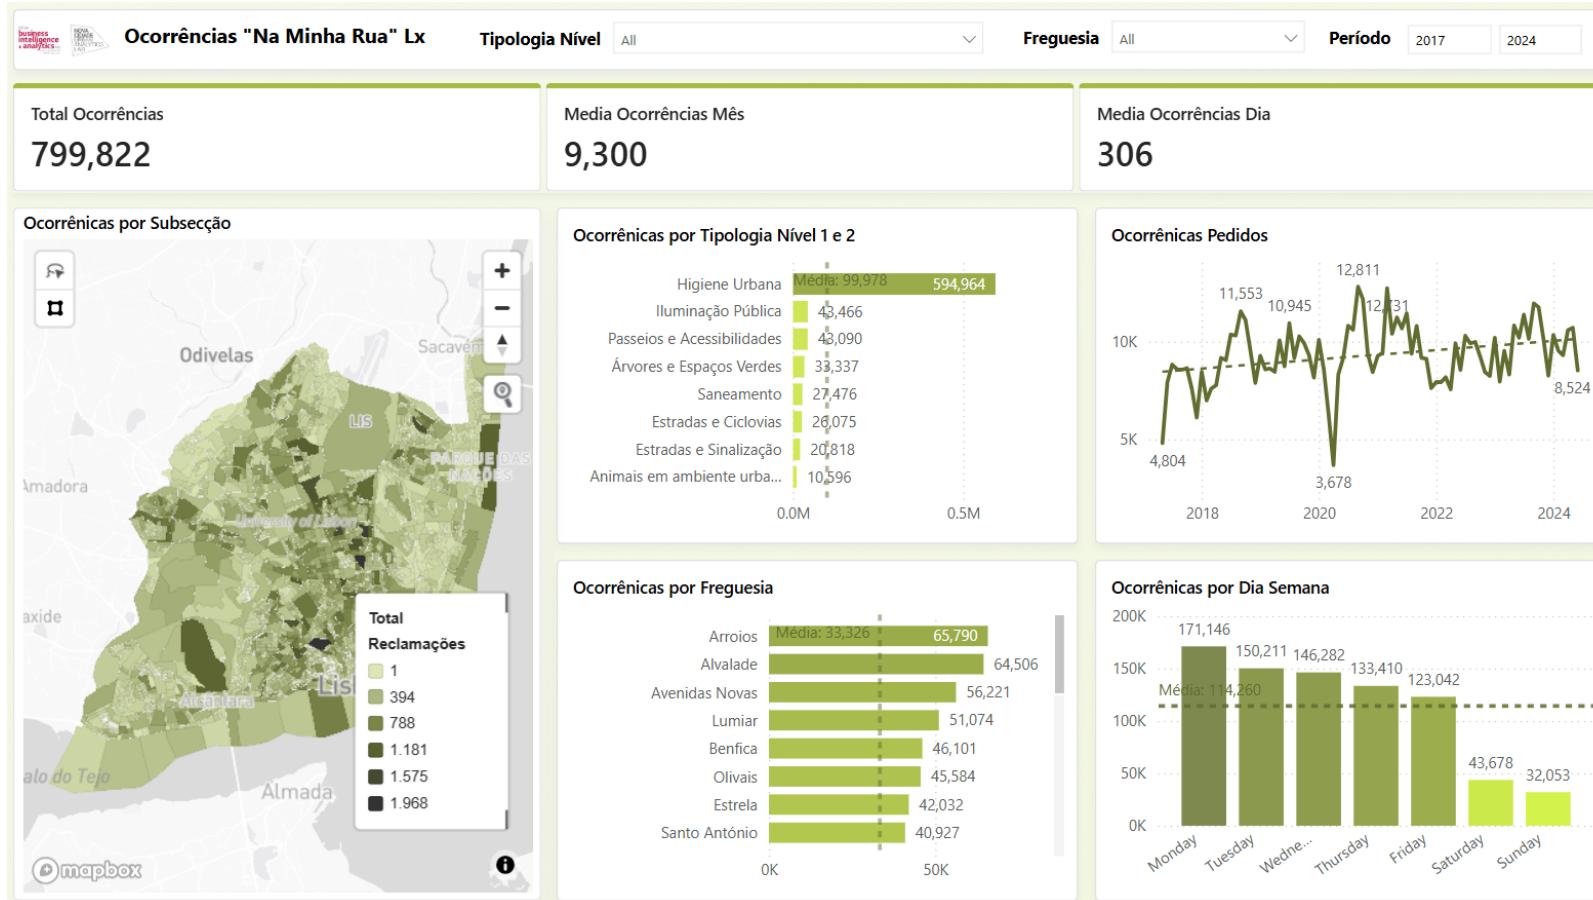
# Example of an applied case (cont.)

Allows for downstream analysis of illegal parking classes in different contexts



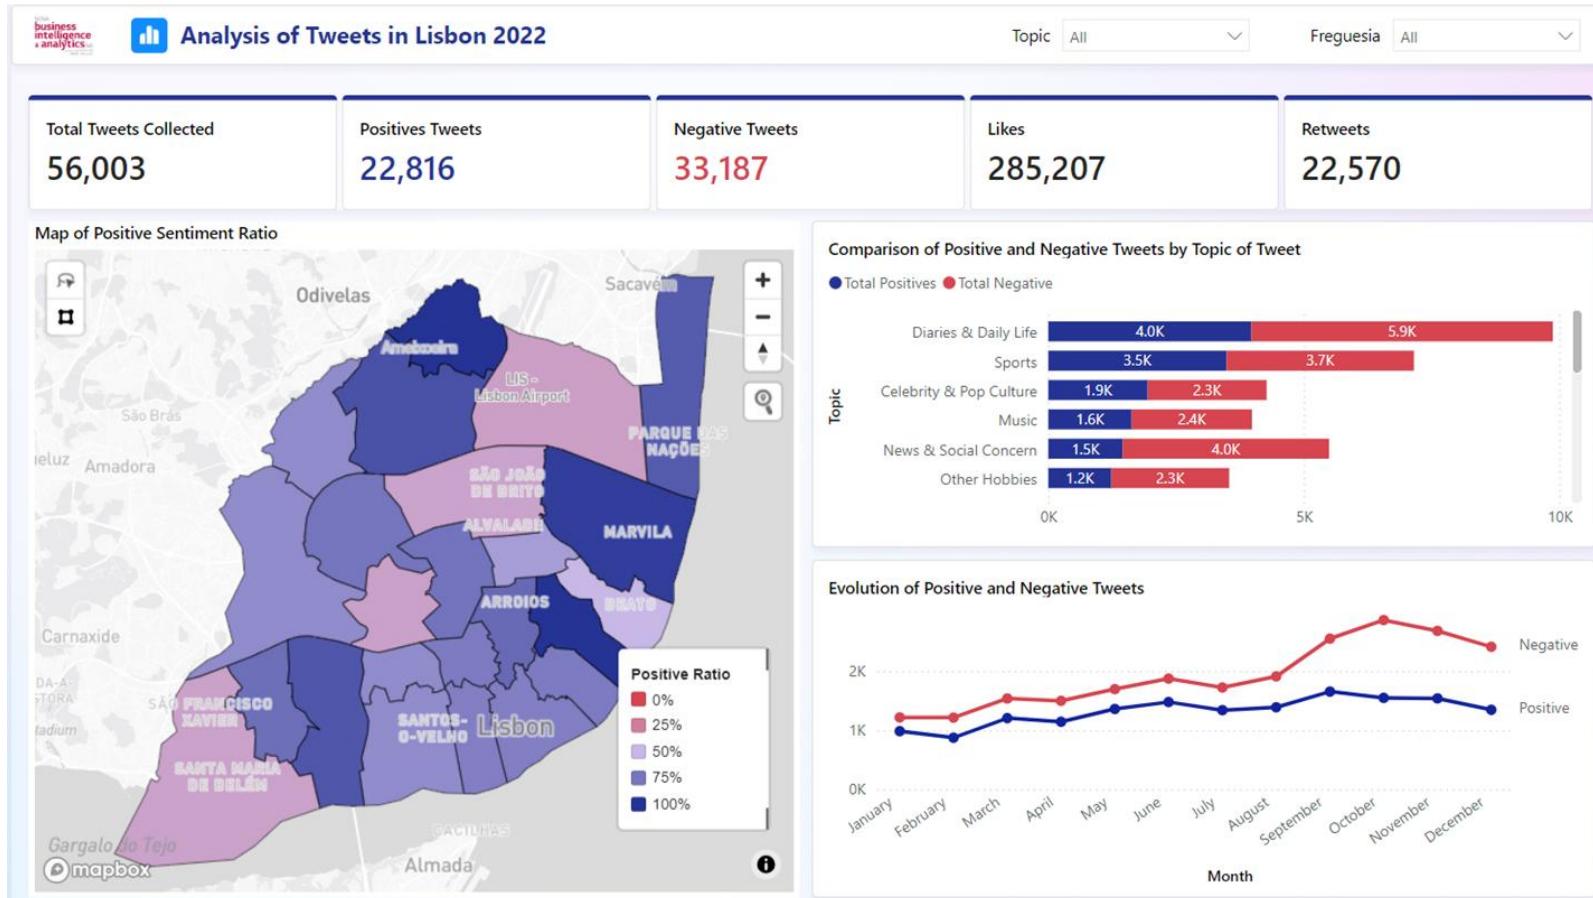
# Example of an applied case

## Public Space Occurrences and Claims



# Example of an applied case

## Analysis of Tweets (X) in Lisbon 2022



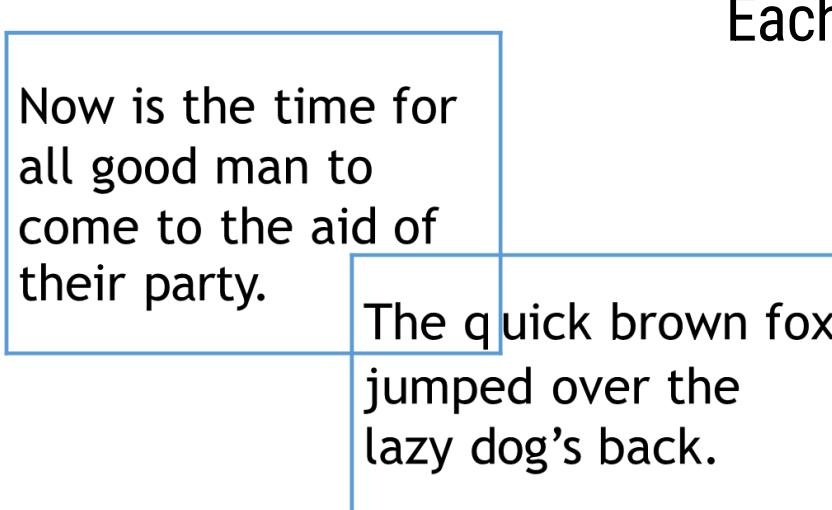
<https://lnkd.in/dRrfDybH>

## 3. Text Preprocessing

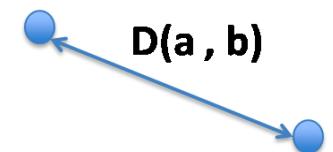
# Text Preprocessing

## Bag-of-Words

Each word is a feature!



$$D(a, b) = \sqrt{\sum_{i=1}^n (b_i - a_i)^2}$$



|      | aid | all | back | brown | come | dog | fox | good | jump | lazy | men | now | over | party | quick | their | Time |
|------|-----|-----|------|-------|------|-----|-----|------|------|------|-----|-----|------|-------|-------|-------|------|
| Doc1 | 0   | 0   | 1    | 1     | 0    | 1   | 1   | 0    | 1    | 1    | 0   | 0   | 1    | 0     | 1     | 0     | 0    |
| Doc2 | 1   | 1   | 0    | 0     | 1    | 0   | 0   | 1    | 0    | 0    | 1   | 1   | 0    | 1     | 0     | 1     | 1    |

# Tokenization

**Tokenization** (or word segmentation) is the process of **identifying tokens/features in text**.

This process is **not trivial**:

- United Kingdom ? Is it two words or one?
- Sexta-feira?
- I am vs I'm?

# Tokenization

## Compounds:

It might be important to treat all compound terms, such as “United Kingdom” as unique features (they share the same semantics)

This does not occur only in names: “soap opera”, “environmentally-friendly”, “fat free”

## Assuming a Bag-of-words model:

“This cake is free fat”

**vs**

“This cake is fat free”

# Punctuation

**Punctuation** can bring several problems to the tokenization.

Typically, we **treat punctuation as individual tokens** but sometimes this can bring problems:

“...” – [., ., .]

“Dr.” – [Dr, .]

“www.google.com” – [www, ., google, ., com]

**vs**

Hey! – [Hey, !]

I love NLP. – [I, love, NLP, .]

# Problems with BoW

1. Curse of Dimensionality

2. No Semantic Relationships

# Curse of Dimensionality

It has the **curse of dimensionality** issue as the **total dimension of the BoW is the vocabulary size**. It can easily over-fit your model.

The remedy is to **use some well-known dimensionality reduction** technique to your input data.

# Curse of Dimensionality

**Bag-of-Words size equals the size of the vocabulary**

Text 1 – “I love Paris”

Text 2 – “ I live in France”

|               | <b>France</b> | <b>I</b> | <b>in</b> | <b>live</b> | <b>love</b> | <b>Paris</b> | <b>(...)</b> | <b>word n-1</b> | <b>word n</b> |
|---------------|---------------|----------|-----------|-------------|-------------|--------------|--------------|-----------------|---------------|
| <b>Text 1</b> | 0             | 1        | 0         | 0           | 1           | 1            | ...          | 0               | 0             |
| <b>Text 2</b> | 1             | 1        | 1         | 1           | 0           | 0            | ...          | 0               | 0             |

# Curse of Dimensionality

Bag of words representation **doesn't consider the semantic relation** between words.

- **Word order is discarded**
- **Similar words are treated as different features** and do not share any information...

**E.g.:** Paris and France should share some information since, for example, they are both locations

# Text Preprocessing



# Eliminating Stop Words

**Some words are not very informative, and their presence only represents more rules and/or more processing.** Thus, the first thing we can do is to get rid of them.

Examples: on, the, a, of, from, and, that, it, in (...)

“Who **is the** main actor **of** Pulp Fiction”

- [who, main, actor, Pulp Fiction]

**Less tokens same meaning!**

“What actor played **the** main role **in** Pulp Fiction”

- [what, main, actor, Pulp Fiction, played, role]

# Lowercasing

Converting everything to lowercase also helps reducing the vocabulary size.  
Examples:

“Comi sopa e comi muito bem” – [Comi, sopa, comi, muito, bem]

“Comi sopa e comi muito bem” – [sopa, comi, muito, bem]

**Sometimes it can go wrong.... Bush (name) != bush (plant)**

# Regular Expressions

In some situations, you may want to normalize dates, numbers, names, etc.

Examples:

"03-04-18" vs. "03/04/18"

"John Fitzgerald Kennedy" vs. "John F. Kennedy" vs. "John Kennedy "

Using **regular expressions** we can **convert dates, names and number from one format to another**.

Or, we can **replace them with a default token**. In industry this is an important step for anonymization!

# Regular Expressions

Text 1 – “I love Paris”

Text 2 – “ I live in France”

|               | <b>France</b> | <b>I</b> | <b>in</b> | <b>live</b> | <b>love</b> | <b>Paris</b> |
|---------------|---------------|----------|-----------|-------------|-------------|--------------|
| <b>Text 1</b> | 0             | 1        | 0         | 0           | 1           | 1            |
| <b>Text 2</b> | 1             | 1        | 1         | 1           | 0           | 0            |

# Regular Expressions

Text 1 – “I love Paris”

Text 2 – “ I live in France”

|        | France | I | in | live | love | Paris |
|--------|--------|---|----|------|------|-------|
| Text 1 | 0      | 1 | 0  | 0    | 1    | 1     |
| Text 2 | 1      | 1 | 1  | 1    | 0    | 0     |

$$D(\text{Text 1}, \text{Text 2}) = \sqrt{(1-0)^2 + (1-1)^2 + (1-0)^2 + (1-0)^2 + (0-1)^2 + (0-1)^2} = 2.24$$

# Regular Expressions

Text 1 – “I love ~~Paris~~” → “I love **LOCATION**”

Text 2 – “ I live in ~~France~~” → “I live in **LOCATION**”

|        | France | I | in | live | love | Paris |
|--------|--------|---|----|------|------|-------|
| Text 1 | 0      | 1 | 0  | 0    | 1    | 1     |
| Text 2 | 1      | 1 | 1  | 1    | 0    | 0     |

# Regular Expressions

Text 1 – “I love LOCATION”

Text 2 – “ I live in LOCATION”

|        | I | in | live | love | LOCATION |
|--------|---|----|------|------|----------|
| Text 1 | 1 | 0  | 0    | 1    | 1        |
| Text 2 | 1 | 1  | 1    | 0    | 1        |

# Regular Expressions

Text 1 – “I love LOCATION”

Text 2 – “ I live in LOCATION”

|        | I | in | live | love | LOCATION |
|--------|---|----|------|------|----------|
| Text 1 | 1 | 0  | 0    | 1    | 1        |
| Text 2 | 1 | 1  | 1    | 0    | 1        |

$$D(\text{Text 1}, \text{Text 2}) = \sqrt{(1-1)^2 + (1-0)^2 + (1-0)^2 + (0-1)^2 + (1-1)^2} = 1.73 (< 2.24)$$

# Stemming & Lemmatization

Transform words in a way that if they represent the same meaning they are captured by the same token. Reduce inflection in language (Inflected Language)

- **Stemming** is the process of removing the last few characters of a given word, to obtain a shorter form, even if that form doesn't have any meaning.  
Ex: **Change, changing, changes, changer** becomes **chang**
- **Lemmatization** is the process of turning words into their root word.  
Ex: **Change, changing, changes, changer** becomes **change**

# Part-of-speech filtering:

Consider only certain type of words (e.g: consider only verbs, nouns, adjectives and adverbs).

|     |      |      |         |     |     |       |
|-----|------|------|---------|-----|-----|-------|
| The | Dog  | Is   | Sitting | On  | The | Table |
| DET | NOUN | VERB | VERB    | ADP | DET | NOUN  |
| X   |      |      |         | X   | X   |       |

# Preprocessing

## 1. Normalization

- Replace links with a special token
- Normalize Dates

## 2. Lowercasing

## 3. Tokenization:

- Compounds
- Punctuation

## 4. Remove Stop-Words

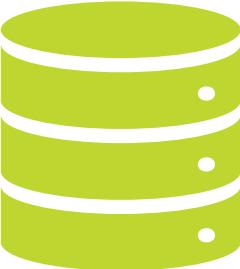
## 5. Stemming and lemmatisation

## 6. POS filtering

Typically, in this order, but everything is optional!!

# Preprocessing

**Corpus**



## **Bag-of-Words**

- Tokenization
- Lowercasing
- Stop words
- Regular Expression



- Keyword Extraction
- Text Similarity
- Document Classification
- Sentiment Analysis
- Topic Modelling
- Information Retrieval
- Q&A
- Text Generation

# Questions and doubts



# Introduction to Text Mining

**Complete these slides by reading the following material:**

<https://web.stanford.edu/~jurafsky/slp3/6.pdf>

pp.1 to pp.11

# Bibliography

- Dan Jurafsky and James H. Martin. **Speech and Language Processing** 3rd ed.  
<https://web.stanford.edu/~jurafsky/slp3/>
- Alammar, J., & Grootendorst, M. (2024). **Hands-On large language models: Language Understanding and Generation**. Sebastopol : O'Reilly Media.  
NOVA IMS Access: <https://search.library.novaims.unl.pt/cgi-bin/koha/opac-detail.pl?biblionumber=97234>

**Thank you!**